



H1 -- cost of paged indirection (decode throughput)



Contiguous beats paged in every cell. Live KV-cache memory differs by ≤ 2 MB between layouts, so the gap is read-path indirection (paged `read()` materialises history via `torch.cat()`), not cache footprint. (2048 $\times$ 16 omitted: 'slow' for both layouts.)