



H2 -- block size is a fragmentation $\leftrightarrow$ latency trade-off, not a single optimum

(a) fragmentation RISES ▲



(b) per-token latency FALLS ▼



Opposing directions are the point: small blocks waste less memory but cost more per-token latency. Throughput-based block-size ranking is **WITHDRAWN** (launch-bound, not reproducible) and is not shown. TPOT is also taxed by the paged read-path copy (a measurement caveat), so panel (b) is an upper bound on the true indirection cost.